

DATA607: Project1

Neil Shah

2/13/2020

#DATA607: Project 1

Introduction:

This project will demonstrate the use of R in reading in the parsing of a text-file, which is then loaded into a dataframe. This is a very common task in data munging.

I collaborated with Jeff Shamp; while our final codes might be different we bounced ideas off each other, and he was very helpful in developing a strategy.

Initializing

Loading a helpful package

```
library('stringr')
```

The text file was saved locally on my computer.

But it can also be found here

'<https://raw.githubusercontent.com/shahneilp/DATA607/master/Project%201/tournamentinfo.txt>
(<https://raw.githubusercontent.com/shahneilp/DATA607/master/Project%201/tournamentinfo.txt>)'

```
filepath <- "tournamentinfo.txt"
```

```
data <- readLines('https://raw.githubusercontent.com/shahneilp/DATA607/master/Project%201/tournamentinfo.txt')
```

Note—I chose to read the file directly into memory; normally this would be the standard due if this was a HUGE file but it isn't and I reasoned that most traditional computers have the RAM to handle it. For big data files I would do it line by line.

Initial Analysis

Opening the text file I see there are two types of rows.

```
data(head)
```

| Pair | Player Name | Total | Round | Round | Round | Round | Round | Round | Round |
|------|---------------------------|-------|-------|-------|-------|-------|-------|-------|-------|
| Num | USCF ID / Rtg (Pre->Post) | Pts | 1 | 2 | 3 | 4 | 5 | 6 | 7 |
| 1 | GARY HUA | 6.0 | W 39 | W 21 | W 18 | W 14 | W 7 | D 12 | D 4 |
| ON | 15445895 / R: 1794 ->1817 | N:2 | W | B | W | B | W | B | W |
| 2 | DAKSHESH DARURI | 6.0 | W 63 | W 58 | L 4 | W 17 | W 16 | W 20 | W 7 |
| MI | 14598900 / R: 1553 ->1663 | N:2 | B | W | B | W | B | W | B |

1. Data containing words with normal white space and punctuation that needs to be extracted.
2. Hyphen separating rows that I can safely eliminate.
3. The pertinent data starts on line 5
4. Player data is every 3rd line

Furthermore:

1. ID are stored as an 8 digit number
2. Total pts are a two digit float
3. States are two capital letter preceded/followed by a white space
4. Names are all capital letters with whitespace—they are preceded by numbers.
5. Pre-scores can be a 3-4 digit number and include a P—they come after a R: and before a ->
6. Each row is 89 characters
7. The player schedule is the opponents rank but can also include games not played.

I will exploit these formatting to create reg-ex to pull said values and create vectors.

Initial RegEx

1. IDs are 8 digit numbers without any white space—this expression in regex is the following:

```
idregex <- "[0-9]{8}"
```

Matching exactly an 8 digit number

2. Total pts are two digit number with floats (decimal) place

```
pointsregex <- "\\d\\.\\d"
```

3. Matching the states—capital letters preceded by a white space on the start of a string with exactly 2 characters.

```
stateregex <- "^\\s+[A-Z]{2}"
```

4. Names are capital strings, at least two characters with white space and preceded by a number and bar |.

```
nameregex <- "[0-9] \\| ([A-Z]+\\s){2,}"
```

5. Scores are 3-4 digit number with possible P string—but are preceded by a R:

```
scoreregex <- "R:\\s*([0-9P]{1,})"
```

Creating Initial vectors

Using the regex expressions I can make column vectors with our data.

```
states <- unlist(str_extract_all(unlist(data), stateregex))
names <- unlist(str_extract_all(unlist(data), nameregex))
points <- unlist(str_extract_all(unlist(data), pointsregex))
ids <- unlist(str_extract_all(unlist(data), idregex))
scores <- unlist(str_extract_all(unlist(data), scoreregex))
```

Post Processing

Let's just check our columns

```
> head(names,2)
[1] "1 | GARY HUA "      "2 | DAKSHESH DARURI "
> head(states,2)
[1] " ON" " MI"
> head(points,2)
[1] "6.0" "6.0"
> head(scores,2)
[1] "R: 1794" "R: 1553"
> head(ids,2)
[1] "15445895" "14598900"
```

Let's remove the whitespace from states, the RL from scores and the number | from names.

For scores—we can just remove the R: portion and keep the number before the P—this is usually a 3,4 digit number. For states we just need to match exactly two capital letters For names—we can just pull the names with spaces and match at least two cases.

Also let's convert the numerical columns to such.

```

> scores <- unlist(str_extract_all(unlist(scores), '[0-9]{3,4}'))
> head(scores,2)
[1] "1794" "1553"

> states <- unlist(str_extract_all(unlist(states), '[A-Z]{2}'))
> head(states,2)
[1] "ON" "MI"

> names <- unlist(str_extract_all(unlist(names), "([A-Z]+\\s){2,}"))
> head(names,2)
[1] "GARY HUA " "DAKSHESH DARURI "

points <- as.numeric(points)
scores <- as.numeric(scores)

```

Data Validation

If we pulled this data correctly they should all have the same entries/length

```

> identical(length(states),length(names),length(points),length(scores))
[1] TRUE

```

Assembling the data frame.

Note: I included ID even though project isn't asking it

```

chessdf <- data.frame(names,states,ids,points,scores)
> head(chessdf)

```

| | names | states | ids | points | scores |
|---|-----------------|--------|----------|--------|--------|
| 1 | GARY HUA | ON | 15445895 | 6.0 | 1794 |
| 2 | DAKSHESH DARURI | MI | 14598900 | 6.0 | 1553 |

Opponent rankings

The opponent played is a series of up to 7 entries, delineated by the rank of an opponent and can include empty values after a D/H. This is tricky since pulling numbers before the | could miss games played. Since we know there is a possibly (max) of 7 games played and 64 players—our end result should be a 64*7 matrix that can include blank/0 values.

Going through the data—the values of the game outcome can be W,L,D,H,U,X,B. Some of these have numbers (meaning a game was played) and some are blank [B,U,D,H,X]. To complicate things the line after each row also contains W,B with blanks.

I noticed that the only data I needed was from each player row so I sliced the data and did an initial regex to find everything between |

```

dataslice <- data[seq(5,length(data),3)]

```

For simplicity I'm going to make matrix/vector called opp to do the calculations.

First to regex anything after a |WDBHXLU

```
opp <- unlist(str_extract_all(unlist(dataslice), "\\|[WDBDXHLU]\\s\\s(.*)"))
```

Replace values of U,H,X,B with 0 to make my life easier—these imply a game not played.

```
opp<-str_replace_all(opp,'H','0')
opp<-str_replace_all(opp,'U','0')
opp<-str_replace_all(opp,'B','0')
opp<-str_replace_all(opp,'X','0')
```

And finally pulling just the numerical values.

```
opp<- unlist(str_extract_all(unlist(opp), "[0-9]{1,2}"))
> head(opp)
[1] "39" "21" "18" "14" "7"  "12"
> length(opp)
[1] 448

opp <- as.numeric(opp)
```

We now have 64*7 or 448 entries.

```
> opp
 [1] 39 21 18 14  7 12  4 63 58  4 17 16 20  7  8 61 25 21 11 13 12 23 28  2 26  5 19  1 45 37
12 13  4 14 17 34 29
 [38] 11 35 10 27 21 57 46 13 11  1  9  2  3 32 14  9 47 28 19 25 18 59  8 26  7 20 16 19 55 31
 6 25 18 38 56  6  7
 [75]  3 34 26 42 33  5 38  0  1  3 36 27  7  5 33  3 32 54 44  8  1 27  5 31 19 16 30 22 54 33
38 10 15  0 39  2 36
[112]  0 48 41 26  2 23 22  5 47  9  1 32 19 38 10 15 10 52 28 18  4  8 40 49 23 41 28  2  9 43
 1 47  3 40 39  6 64
[149] 52 28 15  0 17 40  4 43 20 58 17 37 46 28 47 43 25 60 44 39  9 53  3 24 34 10 47 49 40 17
 4  9 32 11 51 13 46
[186] 37 14  6  0 24  4 22 19 20  8 36 50  6 38 34 52 48  0 52 64 15 55 31 61 50 58 55 64 10 30
50 14 61  8 44 18 51
[223] 26 13 60 12 50 36 13 15 51  6 60 37 29 25 11 52 46 38 56  6 57 52 48 13 57 51 33  0 16 28
 0  5 34 27  0 23 61
[260] 11 35 29 12  0 18 15  1 54 40 16 44 21 24 20 26 39 59 21 56 22 59 17 58 20  0  0  0 12 50
57 60 61 64 56 21 23
[297] 24 63 59 46 55  0 14 32 53 39 24 59  5 51 60 56 63 55 58 35  7 27 50 64 43 23 18 24 21 61
 8 51 25 17 63  0 52
[334]  0 29 35 26 20 63 64 58  0  0 29 42 33 46  0 31 30 27 45 36 57 32 47 33 30 22 19 48 29 35
34  0 25  0 44  0 57
[371]  0 14 39 61  0 15 59 64 62 31 10 30  0 45 43  0 11 35 45  0 40 42  7 36 42 51 35 53  0 31
 2 41 23 49  0 45 41
[408]  0  9 40 43 54 44 33 34 45 42 24  0  0 32  3 54 47 42 30 37 55  0  0  0  0  0  0  2 48 49
43 45  0  0 22 30 31
[445] 49 46 42 54
```

Here we got vector that shows the ranking (player) that each player competed against; 0's imply no opponent and thus not a game played.

Looping over this vector and since chessdf has player scores stored as an index—and keeping mind non-zero indexing.

```
for (r in 1:448) {if (opp[r]==0){opp[r] <-0} else{opp[r]<-as.vector(chessdf$scores)[opp[r]]}}
```

Organizing this into a dataframe

```
opp <-matrix(opp, ncol=7, byrow=TRUE)
opp <-as.data.frame(opp)
> head(opp)
  V1  V2  V3  V4  V5  V6  V7
1 1436 1563 1600 1610 1649 1663 1716
2 1175  917 1716 1629 1604 1595 1649
3 1641  955 1745 1563 1712 1666 1663
4 1363 1507 1553 1579 1655 1564 1794
5 1242  980 1663 1666 1716 1610 1629
6 1399 1602 1712 1438 1365 1552 1563
```

We now have a dataframe with the scores of opponents—non zero values conveniently implies a game was played.

Summing non zero-values

```
> opp <-cbind(opp, Count = rowSums(!opp))
> opp
  V1  V2  V3  V4  V5  V6  V7 Count
1 1436 1563 1600 1610 1649 1663 1716    0
2 1175  917 1716 1629 1604 1595 1649    0
3 1641  955 1745 1563 1712 1666 1663    0
4 1363 1507 1553 1579 1655 1564 1794    0
5 1242  980 1663 1666 1716 1610 1629    0
```

Now simply summing the first 7 elements of the row divided by 7-count, would give us the post rating. I will floor it to keep it as an integer.

We can add this to our chessdf

```
chessdf$postscore <-floor((rowSums(opp[1:7]))/(7-opp$Count)))
```

The Full Dataset

> chessdf

| | names | states | ids | points | scores | postscore |
|----|--------------------------|--------|----------|--------|--------|-----------|
| 1 | GARY HUA | ON | 15445895 | 6.0 | 1794 | 1605 |
| 2 | DAKSHESH DARURI | MI | 14598900 | 6.0 | 1553 | 1469 |
| 3 | ADITYA BAJAJ | MI | 14959604 | 6.0 | 1384 | 1563 |
| 4 | PATRICK H SCHILLING | MI | 12616049 | 5.5 | 1716 | 1573 |
| 5 | HANSHI ZUO | MI | 14601533 | 5.5 | 1655 | 1500 |
| 6 | HANSEN SONG | OH | 15055204 | 5.0 | 1686 | 1518 |
| 7 | GARY DEE SWATHELL | MI | 11146376 | 5.0 | 1649 | 1372 |
| 8 | EZEKIEL HOUGHTON | MI | 15142253 | 5.0 | 1641 | 1468 |
| 9 | STEFANO LEE | ON | 14954524 | 5.0 | 1411 | 1523 |
| 10 | ANVIT RAO | MI | 14150362 | 5.0 | 1365 | 1554 |
| 11 | CAMERON WILLIAM MC LEMAN | MI | 12581589 | 4.5 | 1712 | 1467 |
| 12 | KENNETH J TACK | MI | 12681257 | 4.5 | 1663 | 1506 |
| 13 | TORRANCE HENRY JR | MI | 15082995 | 4.5 | 1666 | 1497 |
| 14 | BRADLEY SHAW | MI | 10131499 | 4.5 | 1610 | 1515 |
| 15 | ZACHARY JAMES HOUGHTON | MI | 15619130 | 4.5 | 1220 | 1483 |
| 16 | MIKE NIKITIN | MI | 10295068 | 4.0 | 1604 | 1385 |
| 17 | RONALD GRZEGORCZYK | MI | 10297702 | 4.0 | 1629 | 1498 |
| 18 | DAVID SUNDEEN | MI | 11342094 | 4.0 | 1600 | 1480 |
| 19 | DIPANKAR ROY | MI | 14862333 | 4.0 | 1564 | 1426 |
| 20 | JASON ZHENG | MI | 14529060 | 4.0 | 1595 | 1410 |
| 21 | DINH DANG BUI | ON | 15495066 | 4.0 | 1563 | 1470 |
| 22 | EUGENE L MCCLURE | MI | 12405534 | 4.0 | 1555 | 1300 |
| 23 | ALAN BUI | ON | 15030142 | 4.0 | 1363 | 1213 |
| 24 | MICHAEL R ALDRICH | MI | 13469010 | 4.0 | 1229 | 1357 |
| 25 | LOREN SCHWIEBERT | MI | 12486656 | 3.5 | 1745 | 1363 |
| 26 | MAX ZHU | ON | 15131520 | 3.5 | 1579 | 1506 |
| 27 | GAURAV GIDWANI | MI | 14476567 | 3.5 | 1552 | 1221 |
| 28 | SOFIA ADINA | MI | 14882954 | 3.5 | 1507 | 1522 |
| 29 | CHIEDOZIE OKORIE | MI | 15323285 | 3.5 | 1602 | 1313 |
| 30 | GEORGE AVERY JONES | ON | 12577178 | 3.5 | 1522 | 1144 |
| 31 | RISHI SHETTY | MI | 15131618 | 3.5 | 1494 | 1259 |
| 32 | JOSHUA PHILIP MATHEWS | ON | 14073750 | 3.5 | 1441 | 1378 |
| 33 | JADE GE | MI | 14691842 | 3.5 | 1449 | 1276 |
| 34 | MICHAEL JEFFERY THOMAS | MI | 15051807 | 3.5 | 1399 | 1375 |
| 35 | JOSHUA DAVID LEE | MI | 14601397 | 3.5 | 1438 | 1149 |
| 36 | SIDDHARTH JHA | MI | 14773163 | 3.5 | 1355 | 1388 |
| 37 | AMIYATOSH PWNANANDAM | MI | 15489571 | 3.5 | 980 | 1384 |
| 38 | BRIAN LIU | MI | 15108523 | 3.0 | 1423 | 1539 |
| 39 | JOEL R HENDON | MI | 12923035 | 3.0 | 1436 | 1429 |
| 40 | FOREST ZHANG | MI | 14892710 | 3.0 | 1348 | 1390 |
| 41 | KYLE WILLIAM MURPHY | MI | 15761443 | 3.0 | 1403 | 1248 |
| 42 | JARED GE | MI | 14462326 | 3.0 | 1332 | 1149 |
| 43 | ROBERT GLEN VASEY | MI | 14101068 | 3.0 | 1283 | 1106 |
| 44 | JUSTIN D SCHILLING | MI | 15323504 | 3.0 | 1199 | 1327 |
| 45 | DEREK YAN | MI | 15372807 | 3.0 | 1242 | 1152 |
| 46 | JACOB ALEXANDER LAVALLEY | MI | 15490981 | 3.0 | 377 | 1357 |
| 47 | ERIC WRIGHT | MI | 12533115 | 2.5 | 1362 | 1392 |
| 48 | DANIEL KHAIN | MI | 14369165 | 2.5 | 1382 | 1355 |
| 49 | MICHAEL J MARTIN | MI | 12531685 | 2.5 | 1291 | 1285 |
| 50 | SHIVAM JHA | MI | 14773178 | 2.5 | 1056 | 1296 |
| 51 | TEJAS AYYAGARI | MI | 15205474 | 2.5 | 1011 | 1356 |

| | | | | | | |
|----|----------------------|----|----------|-----|------|------|
| 52 | ETHAN GUO | MI | 14918803 | 2.5 | 935 | 1494 |
| 53 | JOSE C YBARRA | MI | 12578849 | 2.0 | 1393 | 1345 |
| 54 | LARRY HODGE | MI | 12836773 | 2.0 | 1270 | 1206 |
| 55 | ALEX KONG | MI | 15412571 | 2.0 | 1186 | 1406 |
| 56 | MARISA RICCI | MI | 14679887 | 2.0 | 1153 | 1414 |
| 57 | MICHAEL LU | MI | 15113330 | 2.0 | 1092 | 1363 |
| 58 | VIRAJ MOHILE | MI | 14700365 | 2.0 | 917 | 1391 |
| 59 | SEAN M MC CORMICK | MI | 12841036 | 2.0 | 853 | 1319 |
| 60 | JULIA SHEN | MI | 14579262 | 1.5 | 967 | 1330 |
| 61 | JEZZEL FARKAS | ON | 15771592 | 1.5 | 955 | 1327 |
| 62 | ASHWIN BALAJI | MI | 15219542 | 1.0 | 1530 | 1186 |
| 63 | THOMAS JOSEPH HOSMER | MI | 15057092 | 1.0 | 1175 | 1350 |
| 64 | BEN LI | MI | 15006561 | 1.0 | 1163 | 1263 |

Writing

We can simply now write this dataframe out

```
write.table(chessdf, file = "chessratings.csv", sep = ",", row.names = FALSE)
```

Conclusion

In this Project I successfully took a standard text file and pulled the relevant data into a dataframe and outputted a csv. This, I'm sure is a very common task for data scientists—data munging/process is 80% of the job. While the regex was annoying at first—it's really a powerful tool.

Possible improvements

1. Reading line by line—this would be an optimized solution that would help preventing memory issues for big files.
2. Exploiting the fixed lengths to quickly make dataframes—most of the data was at the same interval, I chose to use regex instead to make this more scalable.
- 3) Better regex—ideally I wouldn't need to post process my data pulls; I'm sure there are more efficient ways to regex.
- 4) Combining code—the use of intermediary dataframes was to make my life easier but I could of combined this all in one

Appendix: Full code


```

data <- readLines('https://raw.githubusercontent.com/shahneilp/DATA607/master/Project%201/tournamentinfo.txt')
idregex <- "[0-9]{8}"
pointsregex <- "\\d\\.\\d"
stateregex <- "^\\s+[A-Z]{2}"
nameregex <- '[0-9] \\| ([A-Z]+\\s){2,}'
scoreregex <- "R:\\s*([0-9P]{1,})"
states <- unlist(str_extract_all(unlist(data), stateregex))
names <- unlist(str_extract_all(unlist(data), nameregex))
points <- unlist(str_extract_all(unlist(data), pointsregex))
ids <- unlist(str_extract_all(unlist(data), idregex))
scores <- unlist(str_extract_all(unlist(data), scoreregex))
scores <- unlist(str_extract_all(unlist(scores), '[0-9]{3,4}'))
states <- unlist(str_extract_all(unlist(states), '[A-Z]{2}'))
names <- unlist(str_extract_all(unlist(names), "([A-Z]+\\s){2,}"))
points <- as.numeric(points)
scores <- as.numeric(scores)
print('Data Validation:')
print(identical(length(states), length(names), length(points), length(scores)))
rank <- seq(1, length(names))
chessdf <- data.frame(names, states, ids, points, scores)
dataslice <- data[seq(5, length(data), 3)]
opp <- unlist(str_extract_all(unlist(dataslice), "\\|[WDBDXHLU]\\s\\s(.*)"))
opp <- str_replace_all(opp, 'H', '0')
opp <- str_replace_all(opp, 'U', '0')
opp <- str_replace_all(opp, 'B', '0')
opp <- str_replace_all(opp, 'X', '0')
opp <- unlist(str_extract_all(unlist(opp), "[0-9]{1,2}"))
opp <- as.numeric(opp)
for (r in 1:448) {if (opp[r]==0){opp[r] <- 0} else{opp[r] <- as.vector(chessdf$scores)[opp[r]]}}
opp <- matrix(opp, ncol=7, byrow=TRUE)
opp <- as.data.frame(opp)
opp <- cbind(opp, Count = rowSums(!opp))
chessdf$postscore <- floor((rowSums(opp[1:7])/(7-opp$Count)))
write.table(chessdf, file = "chessratings.csv", sep = ",", row.names = FALSE)

```